

RetinAI

A Minor Project Report

Submitted in partial fulfilment of requirement of the

Degree of

**BACHELOR OF TECHNOLOGY in COMPUTER
SCIENCE & ENGINEERING**

BY

Aadarsh Sahu EN23IT301001

Udit Maheshwari EN23EL301104

Vigya Gupta EN23IT301110

Under the Guidance of

Prof. Amrata Gupta | Prof. Laxmi Kag



**Department of Computer Science & Engineering
Faculty of Engineering
MEDICAPS UNIVERSITY, INDORE- 453331**

APRIL 2026

Report Approval

The project work “**RetinAI**” is hereby approved as a creditable study of an engineering/computer application subject carried out and presented in a manner satisfactory to warrant its acceptance as prerequisite for the Degree for which it has been submitted.

It is to be understood that by this approval the undersigned do not endorse or approve any statement made, opinion expressed, or conclusion drawn therein; but approve the “Project Report” only for the purpose for which it has been submitted.

Internal Examiner:

Name: Amrata Gupta & Laxmi Kag

Designation: Assistant Professor

Affiliation: Mediacaps University

External Examiner

Name:

Designation:

Affiliation:

Declaration

We hereby declare that the project entitled “**RetinAI**” submitted in partial fulfillment for the award of the degree of Bachelor of Technology in Computer Science & Engineering completed under the supervision of **Prof. Amrata Gupta & Prof. Laxmi Kag**, Faculty of Engineering, Medicaps University Indore is an authentic work.

Further, we declare that the content of this Project work, in full or in parts, have neither been taken from any other source nor have been submitted to any other Institute or University for the award of any degree or diploma.

Aadarsh Sahu

Vigya Gupta

Udit Maheshwari

Certificate

We, **Prof. Amrata Gupta & Prof. Laxmi Kag** certify that the project entitled **RetinAI** submitted in partial fulfillment for the award of the degree of Bachelor of Technology by **Aadarsh Sahu, Vigya Gupta, Udit Maheshwari** is the record carried out by them under our guidance and that the work has not formed the basis of award of any other degree elsewhere.

Prof. Amrata Gupta & Prof. Laxmi Kag

Department of Computer Science & Engineering

Medicaps University, Indore

Prof. (Dr.) Kailash Chandra Bandhu

Head of the Department

Computer Science & Engineering

Medicaps University, Indore

Acknowledgements

I would like to express my deepest gratitude to the Honorable Chancellor, **Shri R. C. Mittal**, who has provided me with every facility to successfully carry out this project, and my profound indebtedness to **Prof. (Dr.) D. K. Patnaik**, Vice Chancellor, Medicaps University, whose unfailing support and enthusiasm has always boosted up my morale. I also thank **Prof. (Dr.) Ratnesh Litoriya**, Associate Dean, Faculty of Engineering, Medicaps University, for giving me a chance to work on this project. I would also like to thank my Head of the Department **Prof. (Dr.) Kailash Chandra Bandhu** for his continuous encouragement for the betterment of the project.

It is their help and support, due to which we became able to complete the design and technical report.

Without their support this report would not have been possible.

Special thanks to our guide, Prof. Amrata Gupta & Prof. Laxmi Kag, for providing technical insights that made this design and technical report possible.

Aadarsh Sahu
Vigya Gupta
Udit Maheshwari

B.Tech. III Year (T)
Department of Computer Science & Engineering
Faculty of Engineering
Medicaps University, Indore

Abstract

RetinAI is an AI-powered system designed for the early detection of retinopathy diseases using retinal image analysis. The project leverages deep learning models, particularly convolutional neural networks (CNNs), to automatically identify and classify abnormalities in fundus images. By training on large datasets of labeled retinal scans, RetinAI is capable of detecting signs of conditions such as diabetic retinopathy with high accuracy and consistency.

The system aims to assist healthcare professionals by providing quick, reliable, and cost-effective screening, especially in regions with limited access to specialized ophthalmic care. RetinAI not only reduces the burden on medical experts but also enables early diagnosis, which is critical for preventing vision loss. The platform integrates a user-friendly interface with efficient data processing and secure storage, making it suitable for both clinical and local deployment environments.

Overall, RetinAI demonstrates the potential of artificial intelligence in improving diagnostic efficiency and accessibility in modern healthcare systems.

Keywords:

Artificial Intelligence, Retinopathy, Convolutional Neural Networks (CNN), Medical Image Processing, Early Disease Detection, Ophthalmology, Predictive Modeling, Clinical Decision Support

Table of Contents

		Page No.
	Report Approval	ii
	Declaration	iii
	Certificate	iv
	Acknowledgement	v
	Abstract	vi
	Table of Contents	vii
Chapter 1	Introduction (Whichever is applicable)	1-8
	1.1 Introduction	1
	1.2 Literature Review	2
	1.3 Objectives	4
	1.4 Significance	5
	1.5 Research Design	6
	1.6 Source of Data	6
	1.7 Chapter Scheme	7
Chapter 2	REQUIREMENTS SPECIFICATION	9-13
	2.1 User Characteristics	9
	2.2 Functional Requirements	9
	2.3 Dependencies	11
	2.4 Performance Requirements	11
	2.5 Hardware Requirements	12
	2.6 Constraints & Assumptions	12

Chapter 3	DESIGN (Whichever is applicable)	14 - 26
	3.1 Algorithm (if Applicable)	14
	3.2 Function Oriented Design for procedural approach	15
	3.3 System Design (Whichever is applicable)	15
	3.3.1 Data Flow Diagrams (Level 0, Level1)	16
	3.3.2 Activity Diagram	19
	3.3.3 Flow Chart	21
	3.3.4 Class Diagram	23
	3.3.5 ER Diagram	24
	3.3.6 Sequence diagram	25
	3.4 Database Design	27
	3.4.1 Logical Database Design	27
	3.4.2 Physical Database Design	28
Chapter 4	Implementation, Testing, and Maintenance	29 - 33
	4.1 Introduction to Languages, IDE's, Tools and Technologies used for Implementation	29
	4.2 Testing Techniques and Test Plans (According to project)	31
	4.3 Installation Instructions	32
	4.4 End User Instructions	33
Chapter 5	Results and Discussions	34 - 37
	5.1 User Interface Representation (of Respective Project)	34
	5.2 Brief Description of Various Modules of the system	35
	5.3 Snapshots of system with brief detail of each	35
	5.4 Back Ends Representation (Database to be used)	36
	5.5 Snapshots of Database Tables with brief description	37

Chapter 6	Summary and Conclusions	38
Chapter 7	Future scope	40
	Appendix	43
	Bibliography	45
	List of References (If any)	46

List of Figures

Figure 3.1 – System Architecture Diagram	11
Figure 3.2 – Data Flow Diagram – Level 0 (Context Diagram)	12
Figure 3.3 – Data Flow Diagram – Level 1 (Detailed)	13
Figure 3.4 – Activity Diagram – Message Translation Workflow	13
Figure 3.5 – Flowchart – System Operation	13
Figure 3.6 – Class Diagram	14
Figure 3.7 – Entity-Relationship (ER) Diagram	14
Figure 3.8 – Sequence Diagram – Message Send and Translate	15
Figure 5.1 – Login Page Interface	20
Figure 5.2 – Registration Page Interface	20
Figure 5.3 – Chat Interface with Translation	21
Figure 5.4 – Language Selection Interface	21
Figure 5.5 – Administrator Dashboard	21

List of Tables

Table 2.1 – Functional Requirements Summary	7
Table 2.2 – Non-Functional Requirements	9
Table 2.3 – Hardware Requirements	10
Table 3.1 – Data Dictionary – User Table	15
Table 3.2 – Data Dictionary – Message Table	16
Table 3.3 – Data Dictionary – Chat Table	16
Table 3.4 – Normalization Summary	16
Table 4.1 – Technology Stack Summary	17
Table 4.2 – Test Cases and Expected Results	18

Abbreviations

Abbreviation	Full Form
API	Application Programming Interface
DFD	Data Flow Diagram
ER	Entity-Relationship
GUI	Graphical User Interface
HTTP/HTTPS	Hypertext Transfer Protocol / Secure
IDE	Integrated Development Environment
JSON	JavaScript Object Notation
JWT	JSON Web Token
NLP	Natural Language Processing
NMT	Neural Machine Translation
OOP	Object-Oriented Programming
ORM	Object-Relational Mapping
SRS	Software Requirements Specification
SDS	Software Design Specification
UI	User Interface
UML	Unified Modeling Language
URL	Uniform Resource Locator
WCAG	Web Content Accessibility Guidelines

Notations & Symbols

Symbol / Notation	Meaning / Usage
→	Data flow direction in DFDs and sequence diagrams
[]	Optional element or array notation in data structures
{ }	Set or group of related elements
FR-n	Functional Requirement identifier (e.g., FR-1, FR-2)
NFR-n	Non-Functional Requirement identifier
PK	Primary Key in database table design
FK	Foreign Key in database table design
1NF / 2NF / 3NF	First / Second / Third Normal Form (database normalization)
VARCHAR(n)	Variable-length character string of max n characters
INT	Integer data type in database
TEXT	Long-form text data type in database
DATETIME	Date and time combined data type
BOOLEAN	True/False binary data type
≤	Less than or equal to (used in performance thresholds)

CHAPTER 1: INTRODUCTION

1.1 Introduction

Diabetes mellitus is one of the most rapidly escalating non-communicable diseases of the twenty-first century. The International Diabetes Federation (IDF) projects that the global diabetic population, standing at approximately 537 million adults in 2021, will exceed 783 million by the year 2045. Among the most devastating and yet preventable complications of long-term uncontrolled diabetes is Diabetic Retinopathy (DR) — a progressive microvascular disease of the retina that represents the leading cause of new-onset blindness in the working-age population worldwide.

Diabetic Retinopathy arises from chronic hyperglycaemia-induced structural and functional damage to the retinal capillary network. In its early non-proliferative stages, the disease manifests as microaneurysms (localized ballooning of capillary walls), dot and blot hemorrhages, and hard exudates (lipid deposits from leaking vessels). As the disease progresses, cotton-wool spots, venous beading, and intraretinal microvascular abnormalities (IRMA) develop. In its most severe proliferative stage, the retina responds to ischemia by triggering pathological neovascularization — the growth of fragile, abnormal blood vessels that are prone to rupture, vitreous hemorrhage, tractional retinal detachment, and ultimately irreversible blindness.

The clinical management of Diabetic Retinopathy is fundamentally dependent on early detection. Epidemiological studies have consistently demonstrated that timely intervention — including laser photocoagulation, intravitreal anti-VEGF injections, or vitreoretinal surgery — can prevent over 90% of severe visual loss in patients identified at an early or moderate stage. This creates an imperative for systematic, population-level retinal screening of all diabetic patients at regular intervals.

However, the traditional screening model presents significant operational challenges. The standard diagnostic protocol requires a trained ophthalmologist to examine high-resolution digital fundus photographs using slit-lamp biomicroscopy or indirect ophthalmoscopy. This process is labour-intensive, time-consuming, highly subjective, and prone to inter-observer variability. In India, the ophthalmologist-to-population ratio stands at approximately 1:100,000, which is far below the

WHO-recommended ratio, creating massive screening backlogs, particularly in Tier-2 and Tier-3 cities and rural areas.

The emergence of Artificial Intelligence, particularly Deep Learning and Convolutional Neural Networks (CNNs), has opened an unprecedented opportunity to automate the analysis of medical images with expert-level performance. In 2016, Gulshan et al. demonstrated that a deep neural network trained on over 128,000 retinal images could achieve diagnostic sensitivity and specificity comparable to board-certified ophthalmologists, marking a pivotal moment for AI in ophthalmology.

RetinAI is developed in this context as a practical, scalable, and clinically safe automated screening platform. By strategically framing the problem as binary classification — Healthy vs. Diseased — rather than the academically common but clinically challenging 5-class grading, RetinAI achieves a simpler, faster, and more reliable triage decision. The project integrates Transfer Learning using the ResNet50 architecture with a custom classification head, deployed as a full-stack web application accessible to medical personnel without requiring deep technical expertise.

1.2 Literature Review

The application of computational techniques to retinal image analysis has a history spanning nearly three decades, evolving from traditional rule-based image processing methods to the current dominance of deep learning architectures. This section provides a critical synthesis of the key milestones and the research gaps that RetinAI seeks to address.

The earliest automated DR detection systems relied on morphological image processing techniques. Sinthanayothin et al. (1999) proposed one of the first automated methods for the detection of hard exudates and hemorrhages using recursive region growing algorithms and Multilayer Perceptrons (MLPs). While foundational, such approaches were highly sensitive to image quality variations, illumination inconsistencies, and the wide morphological heterogeneity of retinal lesions across patient populations.

The introduction of Support Vector Machines (SVMs) and Random Forests for classification improved the robustness of traditional feature-based systems. Gardner et al. demonstrated the utility of neural networks for early DR screening, achieving moderate sensitivity values. However, these systems were fundamentally limited by the quality of hand-crafted features, which required extensive domain expertise and failed to generalize across diverse fundus camera hardware.

The deep learning revolution, initiated by the AlexNet breakthrough at ILSVRC 2012, transformed medical image analysis. Architectures such as VGGNet, GoogLeNet (InceptionV3), and ResNet were rapidly adapted for retinal image grading. The landmark publication by Gulshan et al. in JAMA (2016) trained a deep CNN on 128,175 retinal images from EyePACS and Messidor-2 datasets, achieving an AUC of 0.991 for moderate-or-worse DR detection — a performance level exceeding that of most ophthalmologists. This work provided the scientific foundation upon which RetinAI is built.

He et al. (CVPR 2016) introduced the concept of Residual Learning, wherein identity shortcut connections allow training of extremely deep networks (50, 101, 152 layers) without suffering from the vanishing gradient problem. ResNet50 achieved a top-5 error rate of 5.25% on ImageNet, and its hierarchical feature extraction capabilities — from low-level edge detectors in early layers to high-level semantic detectors in deeper layers — have proven exceptionally transferable to medical imaging domains.

Despite significant progress, critical gaps in the existing literature continue to impede real-world clinical adoption. First, the overwhelming majority of published models target five-class severity grading (No DR, Mild, Moderate, Severe, Proliferative), a formulation that introduces significant prediction ambiguity in the moderate stages — precisely where early intervention is most impactful. This multi-class approach requires far larger and more precisely labeled datasets and generates lower per-class confidence scores. RetinAI strategically simplifies this by adopting binary classification, which aligns with the clinical decision required at a screening stage: refer or do not refer.

Second, most published research models optimize for overall statistical Accuracy as the primary metric. In the inherently imbalanced DR dataset — where healthy retinas constitute approximately 73.5% of the EyePACS data — a trivial classifier that always predicts 'Healthy' can achieve 73.5% accuracy. This optimization leads to unacceptably high False Negative rates, which in a clinical context means missing diseased patients. RetinAI explicitly prioritizes Recall (Sensitivity) over Accuracy through class-weighted loss functions, addressing this fundamental clinical safety gap.

Third, most research prototypes are computationally prohibitive for deployment in resource-constrained environments, typically requiring high-memory GPU servers for inference. RetinAI is specifically architected for lightweight inference on standard laptop hardware (8GB RAM, Intel i5), making it viable for deployment in small clinics and primary health centers. Fourth, existing tools

lack a clinician-centric user interface, presenting raw probability vectors or command-line outputs that are meaningless to medical practitioners. RetinAI encapsulates its complex inference engine within a clean, minimalist web dashboard designed for non-technical medical users.

1.3 Objectives

The RetinAI project is governed by the following set of clearly defined technical and clinical objectives:

1. To develop a robust binary classification engine capable of accurately distinguishing between retinal fundus images exhibiting signs of Diabetic Retinopathy (Class 1 – Diseased) and clinically normal retinas (Class 0 – Healthy).
2. To leverage Transfer Learning using the ResNet50 architecture pre-trained on ImageNet to extract rich, hierarchical vascular features from retinal images without requiring the computational overhead of training from scratch.
3. To design and append a custom CNN classification head to the ResNet50 backbone, incorporating Global Average Pooling, Dense layers, Dropout regularization, and a Sigmoid activation output neuron for binary probabilistic classification.
4. To implement a comprehensive data preprocessing and augmentation pipeline using OpenCV and Keras ImageDataGenerator to standardize image dimensions (224×224 pixels), normalize pixel intensities, and artificially expand the training dataset.
5. To prioritize and mathematically enforce clinical Recall (Sensitivity) $\geq 90\%$ through class-weighted Binary Cross-Entropy loss functions, ensuring the minimization of clinically dangerous False Negative predictions.
6. To perform iterative fine-tuning of the ResNet50 backbone layers at a low learning rate following initial classification head training, enabling domain-specific specialization of generic ImageNet features to retinal anatomy.
7. To design and deploy a user-friendly, doctor-centric three-tier web application (Frontend, Backend, Model) providing a seamless image upload, analysis, and result visualization workflow accessible to non-technical medical personnel.
8. To implement a lightweight PostgreSQL database for the anonymous logging of diagnostic outcomes, confidence scores, and system performance metrics in compliance with medical data privacy standards.
9. To ensure end-to-end inference latency of under 5 seconds on the specified development hardware, validating the system's suitability for real-time clinical screening workflows.

1.4 Significance of the Study

The significance of RetinAI extends across multiple dimensions — clinical, social, and technological.

From a clinical perspective, the system directly addresses the global ophthalmologic care gap. By providing an automated, high-sensitivity triage tool, RetinAI enables primary care physicians and general practitioners — who may lack specialized ophthalmic training — to conduct preliminary DR screening. This is particularly transformative in rural and semi-urban regions of India and developing nations, where the density of ophthalmologists remains critically low. A patient flagged as 'Diseased' by RetinAI can be promptly referred to a specialist, while 'Healthy' patients can be safely scheduled for routine follow-up, dramatically improving the efficiency of specialist time allocation.

From a patient safety perspective, the system's aggressive optimization for Recall ensures that virtually no diseased patient is sent home without a referral. In preventive medicine, the principle of 'primum non nocere' (first, do no harm) demands that the cost of missing a sick patient must be considered far greater than the cost of a false alarm. RetinAI's loss function architecture directly encodes this clinical ethic into its mathematical optimization objective.

From a technological perspective, RetinAI demonstrates the viability of deploying sophisticated deep learning models on commodity hardware without cloud dependency. This architectural decision has profound implications for healthcare systems in regions with unreliable internet connectivity, ensuring that the diagnostic tool remains operational as a standalone local application.

From a social and economic perspective, blindness due to undetected DR imposes enormous costs on individuals, families, and healthcare systems. Automated screening at scale, as enabled by RetinAI, offers a cost-effective intervention that can prevent these social and economic burdens.

1.5 Research Design

The research methodology underlying RetinAI follows a structured, empirical design comprising four distinct phases: data acquisition and preparation, model architecture design, training and optimization, and deployment and evaluation.

The research paradigm is applied and experimental in nature. The independent variable is the model architecture configuration (backbone depth, classification head design, loss function weighting), and the dependent variables are the resulting evaluation metrics — primarily Recall and F1-Score. The study employs a quantitative evaluation framework using a held-out test set to measure model performance.

The dataset used is the publicly available EyePACS retinal fundus image dataset from Kaggle, providing thousands of labelled images representing both healthy and diseased retinas across all four stages of DR severity. The binary relabeling strategy — collapsing Grades 1-4 into a single Diseased class — constitutes a deliberate research design decision validated by clinical literature, which identifies any measurable retinopathy as sufficient grounds for specialist referral.

Model training is conducted on Kaggle Notebook GPU environments (NVIDIA Tesla T4) to handle the computational demands of ResNet50 fine-tuning on large image batches, while inference is validated on local hardware to confirm real-time performance on the target deployment platform.

1.6 Source of Data

The primary data source for this project is the EyePACS Diabetic Retinopathy Detection Dataset, publicly hosted on the Kaggle platform (<https://www.kaggle.com/dreamer07/eyepacs>). EyePACS is a tele-ophthalmology service provider based in the United States that has collected one of the largest publicly available retinal fundus image databases in existence.

The dataset consists of high-resolution colour fundus photographs, captured using a variety of fundus camera hardware across multiple clinical sites. Each image is labelled by a trained clinician on a scale of 0 to 4 according to the International Clinical Diabetic Retinopathy Disease Severity Scale: 0 (No apparent retinopathy), 1 (Mild NPDR), 2 (Moderate NPDR), 3 (Severe NPDR), and 4 (Proliferative DR). The dataset exhibits significant class imbalance, with the majority of images belonging to Class 0 (approximately 73.5%).

For the purposes of RetinAI, the five-class labels are collapsed into a binary representation: Class 0 (Healthy) retains all Grade 0 images, while Class 1 (Diseased) aggregates all images with Grade 1 through Grade 4 labels. This relabeling is performed programmatically using a Pandas-based Python script during the data preparation phase.

Secondary data sources include the pre-trained ImageNet weights of the ResNet50 model, accessed through the Keras Applications module (`tensorflow.keras.applications.ResNet50`).

1.7 Chapter Scheme

The report is organized into seven chapters as follows:

- Chapter 1 – Introduction: Presents the background of Diabetic Retinopathy, the motivation for automated screening, a review of relevant literature, the project objectives, significance, research design, and data sources.
- Chapter 2 – Requirements Specification: Provides a detailed Software Requirements Specification (SRS) document including user characteristics, functional requirements (in REQ-X format), non-functional requirements, hardware/software dependencies, performance benchmarks, and design constraints.
- Chapter 3 – System Design: Covers the complete system design model including the algorithmic pipeline, Data Flow Diagrams (Level 0 and Level 1), Activity Diagram, Flowchart of the AI inference module, Class Diagram, Entity-Relationship Diagram, Sequence Diagram, and Database Design with normalization analysis.
- Chapter 4 – Implementation: Details the technology stack, implementation methodology, installation and configuration procedures, and user operating instructions.
- Chapter 5 – Results and Discussion: Presents the user interface description, module breakdown, output results, database structure, and a discussion of model performance metrics.
- Chapter 6 – Conclusion: Summarizes the findings and achievements of the project.
- Chapter 7 – Future Scope: Identifies potential enhancements and extensions for future development cycles.
- Appendix & Bibliography: Contains the glossary of technical terms and IEEE-formatted references.

CHAPTER 2: REQUIREMENTS SPECIFICATION

2.1 User Characteristics

RetinAI is designed to serve two distinct user categories, each with unique technical backgrounds and operational requirements:

2.1.1 Medical Practitioners (Primary Users)

The primary users of RetinAI are medical doctors, ophthalmologists, ophthalmic technicians, and trained lab assistants working in clinical environments ranging from large hospital ophthalmology departments to smaller primary health centres and screening camps. These users possess deep domain expertise in ocular medicine but are not expected to have any programming or data science background. Their interaction with the system must therefore be as frictionless and intuitive as possible, requiring no more than basic computer literacy. The interface must present diagnostic results in unambiguous clinical language — specifically, binary Yes/No referral decisions paired with a confidence percentage — rather than raw probability vectors or statistical scores.

2.1.2 Administrators (Secondary Users)

System administrators are IT personnel responsible for the maintenance, configuration, and updating of the RetinAI platform. Their responsibilities include managing the Flask/FastAPI server infrastructure, administering the user authentication database, deploying updated model weight files as improved versions become available, and monitoring system performance logs. Administrators require technical proficiency in Python environments, web server configuration, and basic database administration.

2.2 Functional Requirements

The following functional requirements define the mandatory behavioural specifications of the RetinAI system, identified using the REQ-X convention for traceability:

2.2.1 User Authentication Module

- REQ-1: The system shall provide a secure login interface requiring a valid username and an encrypted password for all medical practitioner users.
- REQ-2: The system shall authenticate user credentials against an encrypted database entry and grant access exclusively to verified users.

- REQ-3: The system shall implement session management to automatically log out inactive sessions after a configurable timeout period.
- REQ-4: The administrator shall be able to register new user accounts, reset passwords, and deactivate accounts through an admin panel.

2.2.2 Image Upload and Validation Module

- REQ-5: The system shall provide a drag-and-drop upload interface as well as a traditional file browser dialog for retinal fundus image ingestion.
- REQ-6: The system shall validate the uploaded file signature and reject any file that is not a valid JPG or PNG image format, displaying an 'Unsupported Format' error message.
- REQ-7: The system shall enforce an upper file size limit of 20 MB per uploaded image to prevent server memory overload.
- REQ-8: The system shall generate and display a thumbnail preview of the uploaded fundus image prior to analysis initiation.

2.2.3 Preprocessing Module

- REQ-9: The system shall automatically crop black circular borders from uploaded fundus images using OpenCV contour detection algorithms.
- REQ-10: The system shall resize all uploaded images to a standardized 224×224×3 pixel tensor format while preserving as much of the retinal disc as possible.
- REQ-11: The system shall normalize pixel intensity values to the range [0, 1] using ImageNet channel-wise mean and standard deviation values for optimal ResNet50 alignment.

2.2.4 AI Diagnosis Engine

- REQ-12: The system shall execute a complete forward inference pass through the loaded ResNet50-based model and return a binary classification label (Healthy/Diseased) for each uploaded image.
- REQ-13: The system shall output a numerical confidence score between 0% and 100% corresponding to the probability of the Diseased class.
- REQ-14: The system shall trigger a 'Referral Needed – Possible DR Detected' alert when the Diseased class probability exceeds the clinical decision threshold (default: 0.45).
- REQ-15: The system shall complete the end-to-end inference pipeline within a maximum of 5 seconds on the specified development hardware.

2.2.5 Result Display and History Module

- REQ-16: The system shall display the diagnostic result with colour-coded visual feedback: Emerald Green for Healthy and Crimson Red for Diseased.
- REQ-17: The system shall display a probabilistic confidence gauge alongside the classification label to provide clinical context.

- REQ-18: The system shall maintain an anonymized diagnostic history log accessible from the dashboard, displaying past prediction timestamps and confidence scores without storing patient PII.

2.3 Non-Functional Requirements

2.3.1 Performance Requirements

- The end-to-end inference time, measured from image upload completion to result display, shall not exceed 5 seconds on a machine with an Intel Core i5 processor and 8 GB of RAM.
- The system backend shall be capable of handling a minimum of 5 concurrent classification requests from separate browser sessions without significant performance degradation or server crash.
- The system shall achieve a minimum Recall (Sensitivity) of 90% on the held-out test set to satisfy clinical safety requirements.

2.3.2 Security Requirements

- All user passwords shall be stored in hashed form using a cryptographic algorithm such as bcrypt and shall never be stored in plaintext.
- No patient-identifiable information (PII) shall be stored within the diagnostic log database, ensuring compliance with patient privacy regulations.
- The system shall operate exclusively over HTTPS in a production deployment environment to ensure encrypted data transmission.

2.3.3 Usability Requirements

- A first-time user with basic computer literacy shall be able to complete a full diagnostic cycle — from login to result viewing — within 60 seconds without consulting any documentation.
- All error messages displayed by the system shall be written in plain, non-technical language comprehensible to medical professionals without a computing background.

2.3.4 Reliability Requirements

- The system shall maintain an operational uptime of 99% during designated clinic hours.
- The system shall handle malformed or corrupted image files gracefully, displaying an appropriate error message without causing a server crash or exception propagation to the user interface.

2.3.5 Maintainability Requirements

- The system architecture shall be modular, such that the AI model component can be independently replaced (e.g., upgrading from ResNet50 to ResNet101) without requiring modifications to the frontend or authentication modules.

- All backend code shall be documented with inline comments and a comprehensive README file to facilitate future maintenance and enhancement.

2.4 Hardware and Software Dependencies

2.4.1 Hardware Requirements

Component	Minimum Specification	Recommended Specification
Processor	Intel Core i3 (8th Gen)	Intel Core i5 / i7 (10th Gen+)
RAM	4 GB	8 GB or more
Storage	10 GB free space	50 GB SSD
GPU (Training)	Not required (CPU only)	NVIDIA GPU (Kaggle/Colab for training)
Display	1366×768	1920×1080 (Full HD)
Internet	Required for cloud training	Required for initial setup only

2.4.2 Software Requirements

Category	Technology	Version
Operating System	Windows 10/11, Ubuntu 20.04+	Latest Stable
Programming Language	Python	3.9+
Deep Learning Framework	TensorFlow / Keras	2.x
Web Framework	Flask / FastAPI	2.x / 0.100+
Image Processing	OpenCV, PIL (Pillow)	4.x / 9.x
Data Processing	NumPy, Pandas	Latest
Frontend	HTML5, CSS3, JavaScript, React	Tailwind 3.x
Database	SQLite	3.x
Model Serialization	Keras HDF5 (.h5 format)	TF 2.x compatible

2.5 Constraints and Assumptions

2.5.1 Constraints

- **Memory Constraint:** The 8 GB RAM limitation on the target deployment machine necessitates the use of memory-efficient Python data generators (tf.data.Dataset API) and memory-mapped file handling during batch inference. Full dataset retraining on local hardware is not feasible and shall be conducted exclusively on GPU-accelerated cloud environments (Kaggle Notebooks, Google Colab).
- **Data Format Constraint:** The system exclusively accepts JPG and PNG image formats, consistent with standard fundus camera export formats. DICOM format support is not included in the current version scope.
- **Schedule Constraint:** The complete system, from model training to deployed web application, must reach a functional Beta state within the 16-week academic semester timeline.
- **Regulatory Constraint:** RetinAI is classified as a screening and triage assistant and does not constitute a medical diagnostic device. It shall not be used as a standalone clinical decision-making tool in the absence of a qualified ophthalmologist.

2.5.2 Assumptions

- Input images are assumed to be captured using standard clinical fundus cameras (e.g., Topcon, Zeiss, Canon) under adequate illumination conditions. Images captured under severely sub-optimal lighting conditions may yield reduced classification accuracy.
- The system assumes that users have basic computer literacy and access to a modern web browser (Chrome, Firefox, Edge) with JavaScript enabled.
- The Keras Applications library and pre-trained ImageNet weights are assumed to be accessible during the initial model setup phase. Offline deployment may require local caching of these weights.

CHAPTER 3: SYSTEM DESIGN

3.1 Algorithm: Diabetic Retinopathy Detection Pipeline

The core algorithmic pipeline of RetinAI can be formally described as a sequence of deterministic preprocessing transformations followed by a probabilistic deep neural network inference step. The complete pipeline operates as follows:

Step 1 – Image Acquisition and Validation

The system receives an uploaded image file from the web interface. The file is first validated for format integrity (JPG/PNG signature verification). If validation fails, a user-facing error message is generated and the pipeline terminates. Upon successful validation, the image is loaded into memory as a PIL Image object.

Step 2 – Preprocessing

The raw fundus image undergoes a multi-stage preprocessing sequence. First, black circular border regions — which constitute non-informative pixel space from the fundus camera optics — are identified and cropped using contour-based algorithms in OpenCV. Second, the cropped image is resized to a fixed spatial dimension of 224×224 pixels, the mandatory input size for ResNet50. This resizing is performed using bilinear interpolation to minimize aliasing artefacts. Third, the pixel values, originally in the range [0, 255], are normalized by scaling to [0, 1] and subsequently adjusted using the ImageNet channel-wise mean values ($\mu = [0.485, 0.456, 0.406]$ and $\sigma = [0.229, 0.224, 0.225]$ for R, G, B channels respectively), ensuring mathematical alignment with the pre-trained ResNet50 weight distributions. Finally, the resulting tensor is expanded to a batch dimension of shape (1, 224, 224, 3) for model compatibility.

Step 3 – Feature Extraction via ResNet50 Backbone

The normalized (1, 224, 224, 3) tensor is passed through the frozen ResNet50 convolutional backbone. ResNet50 comprises 48 convolutional layers, 1 MaxPooling layer, and 1 Average Pooling layer, organized into four residual stage blocks. Each residual block employs identity shortcut connections that allow the network to learn residual functions $F(x) = H(x) - x$ rather than directly learning the target mapping $H(x)$, which eliminates the vanishing gradient problem and enables effective training of very deep networks. The output of the ResNet50 backbone is a feature map of spatial dimensions 7×7 with 2048 channels, representing a rich, hierarchical abstraction of the retinal image's vascular texture and morphological features.

Step 4 – Custom Classification Head

The (7, 7, 2048) feature map is processed by the custom CNN classification head. A Global Average Pooling (GAP) layer computes the spatial average of each of the 2048 feature maps, producing a compact 1D vector of 2048 values. This approach is superior to Flatten followed by fully connected layers because it dramatically reduces the parameter count, incorporates structural regularization, and enhances translation invariance. The 2048-dimensional vector is then passed through a Dense layer of 512 neurons with ReLU activation, followed by a Dropout layer (rate = 0.5) to prevent overfitting by stochastically zeroing 50% of neurons during training. A second Dense layer of 128 neurons with ReLU activation and Dropout (rate = 0.3) follows. Finally, a single output neuron with Sigmoid activation maps the 128-dimensional vector to a scalar probability value $P(\text{Disease})$ in the range $[0, 1]$.

Step 5 – Decision Logic

The scalar output $P(\text{Disease})$ is compared against a clinical decision threshold θ . If $P(\text{Disease}) > \theta$, the image is classified as Class 1 (Diseased); otherwise it is classified as Class 0 (Healthy). The default threshold is set to $\theta = 0.45$, which is deliberately set below the neutral 0.5 to further prioritize Recall over Precision, increasing the sensitivity of the triage system. This threshold value is determined empirically during the validation phase and can be adjusted by the administrator post-deployment.

Step 6 – Reporting

The binary prediction label and the confidence score ($P(\text{Disease}) \times 100\%$) are serialized into a JSON response by the Flask backend and transmitted to the web frontend. The frontend dynamically renders the result with appropriate colour coding and confidence visualization.

3.3 System Design Models

3.3.1(a) Data Flow Diagram – Level 0 (Context Diagram)

The Level 0 DFD represents the highest-level abstraction of the RetinAI system. At this level, the system is viewed as a single processing entity (a black box) with its external interactions defined.

External Entities: (1) Medical Practitioner – the primary human actor who initiates the diagnostic process by uploading fundus images and receives binary diagnostic results. (2) Administrator – the system maintainer who manages the user database and updates the AI model. (3) Fundus Camera / File System – the external data source from which retinal images are obtained.

Data Flows INTO the system: Retinal Fundus Image (from Practitioner), Login Credentials (from Practitioner), Model Update File (from Administrator), System Configuration Parameters (from Administrator).

Data Flows OUT OF the system: Binary Classification Result (to Practitioner), Confidence Score (to Practitioner), Referral Alert (to Practitioner), Diagnostic History Log (to Administrator), System Performance Reports (to Administrator).

The Level 0 DFD thus captures the boundary of the system and the nature of its interactions with the external world, establishing the scope within which all subsequent design decisions are made.



3.2.1(b) Data Flow Diagram – Level 1 (Decomposition)


The Level 1 DFD decomposes the monolithic RetinAI system process into its four major functional sub-processes, each corresponding to a distinct operational stage:

Process 1.0 – Authentication: This process receives login credentials from the Medical Practitioner and queries the User Data Store (D1) to verify the username-password combination. Upon successful authentication, a session token is generated and returned to the user's browser, granting access to the main dashboard.

Process 2.0 – Image Upload and Preprocessing: This process receives the raw retinal fundus image from the authenticated practitioner. It validates the file format (JPG/PNG), loads the image using PIL, executes the multi-stage preprocessing pipeline (border cropping, resizing to 224×224, normalization), and produces a standardized float32 tensor that is stored temporarily in server memory (D2 – Temp Tensor Store).

Process 3.0 – AI Inference Engine: This process retrieves the preprocessed tensor from D2 and executes the forward inference pass through the loaded ResNet50 Keras model (D3 – Model Weight Store). The ResNet50 backbone extracts the hierarchical feature map, and the custom CNN head generates the final Sigmoid probability score $P(\text{Disease})$. The raw probability is converted to a binary label using the clinical threshold logic.

Process 4.0 – Result Generation and Logging: This process formats the binary label, confidence percentage, and referral recommendation into a structured JSON response. It dispatches this response to the frontend for display and simultaneously writes an anonymized record (timestamp, confidence score, prediction label) to the Diagnostic Log Data Store (D4) in SQLite.




3.3.2 Activity Diagram

The Activity Diagram models the procedural workflow of a single diagnostic session in RetinAI, capturing the sequential and conditional flow of actions from user interaction through system processing to result display.

The workflow initiates with the user logging into the system. A conditional branch checks authentication status: if authentication fails, the system presents an error message and redirects to the login screen. Upon successful authentication, the user is presented with the Analysis Centre dashboard. The user uploads a retinal fundus image through the drag-and-drop zone or file browser. The system immediately validates the file format — if the file fails format validation, an 'Unsupported Format' error is displayed and the upload zone is reset. If validation succeeds, the system generates a thumbnail preview and awaits the user's click on 'Run Analysis.'


Upon clicking 'Run Analysis,' the system executes the preprocessing pipeline (cropping, resizing, normalization) followed by the AI inference pipeline (ResNet50 feature extraction, CNN head classification, Sigmoid decision). A second conditional branch evaluates $P(\text{Disease})$ against the clinical threshold θ : if $P(\text{Disease}) > \theta$, the system renders the result as 'Diseased – Referral Recommended' in Crimson Red; if $P(\text{Disease}) \leq \theta$, the result is rendered as 'Healthy – No Referral Required' in Emerald Green. In both cases, the confidence score is displayed and the diagnostic record is anonymously logged to the database. The workflow concludes when the result is displayed to the user.



3.3.3 Flowchart of the AI Inference Module

The detailed algorithmic flowchart of the AI Inference Module describes the internal logic of the ResNet50 classification pipeline with explicit decision points:

START → Input: Receive raw fundus image file → Decision: Is file JPG or PNG? → [No] → Display error, STOP → [Yes] → Load image using PIL → Apply OpenCV border crop → Resize to 224×224 pixels → Normalize pixel values using ImageNet statistics → Expand tensor to batch shape (1, 224, 224, 3) → Pass tensor through ResNet50 backbone (48 Conv layers + residual connections) → Output: Feature Map (7, 7, 2048) → Apply Global Average Pooling → Output: 2048-D feature vector → Pass through Dense(512, ReLU) + Dropout(0.5) → Pass through Dense(128, ReLU) + Dropout(0.3) → Apply Dense(1, Sigmoid) → Output: P(Disease) in [0, 1] → Decision: P(Disease) > 0.45? → [Yes] → Label = 'Diseased', Alert = 'Referral Needed' → [No] → Label = 'Healthy', Alert = 'No Referral' → Convert P to Confidence% → Log to Database → Return JSON Response → Display Result on Frontend → END.



3.3.4 Class Diagram

The Class Diagram illustrates the static object-oriented structure of the RetinAI software system, defining the core classes and their relationships.

User Class: Attributes include userID (int, PK), username (string), passwordHash (string), role (enum: Doctor/Admin). Methods include login(), logout(), viewHistory().

Image Class: Attributes include imageID (int, PK), filePath (string), format (string), uploadTimestamp (datetime), uploadedByUserID (int, FK). Methods include validate(), generateThumbnail().

Preprocessor Class: Attributes include targetSize (tuple: 224, 224), normMean (list), normStd (list). Methods include cropBorders(image), resize(image), normalize(image), augment(image). This class is an independent utility with no association to User.


ResNet50Classifier Class: Attributes include modelPath (string), threshold (float), backbone (KerasModel). Methods include loadModel(), extractFeatures(tensor), predict(tensor). This class is associated with Preprocessor (uses its output as input) and with the CustomCNNHead class.

CustomCNNHead Class: Attributes include numClasses (int: 1), dropoutRate1 (float), dropoutRate2 (float). Methods include classify(featureVector). This class is a component of ResNet50Classifier.

Result Class: Attributes include resultID (int, PK), imageID (int, FK), label (string), confidence (float). Methods include generateReport(), formatJSON().

DBManager Class: Attributes include connectionString (string). Methods include storeLogs(result), retrieveLogs(userID), manageUsers().

The primary association is: User uploads Image → Image is processed by Preprocessor → ResNet50Classifier uses Preprocessor output → ResNet50Classifier contains CustomCNNHead → ResNet50Classifier generates Result → DBManager stores Result.



3.3.5 Entity-Relationship Diagram


The Entity-Relationship (ER) Diagram models the logical structure of the RetinAI database, defining the entities, their attributes, and the relationships between them.

Entities and Attributes:

- **USER:** user_id (PK, INT), username (VARCHAR, UNIQUE, NOT NULL), password_hash (VARCHAR, NOT NULL), role (VARCHAR: 'doctor'/'admin', NOT NULL), email (VARCHAR), created_at (DATETIME).
- **RETINAL_IMAGE:** image_id (PK, INT), filename (VARCHAR, NOT NULL), file_format (VARCHAR), image_size_kb (FLOAT), upload_timestamp (DATETIME, NOT NULL), uploaded_by (FK → USER.user_id).
- **PREDICTION:** prediction_id (PK, INT), image_id (FK → RETINAL_IMAGE.image_id, UNIQUE), label (VARCHAR: 'Healthy'/'Diseased'), confidence_score (FLOAT, 0-1), threshold_used (FLOAT), prediction_timestamp (DATETIME), model_version (VARCHAR).
- **DIAGNOSTIC_LOG:** log_id (PK, INT), prediction_id (FK → PREDICTION.prediction_id), log_timestamp (DATETIME), system_latency_ms (INT), notes (TEXT).

Relationships:

- **USER to RETINAL_IMAGE:** One-to-Many (one user can upload many images; each image is associated with exactly one user).
- **RETINAL_IMAGE to PREDICTION:** One-to-One (each image produces exactly one prediction result; each prediction belongs to exactly one image).
- **PREDICTION to DIAGNOSTIC_LOG:** One-to-One (each prediction generates exactly one log entry).




3.3.6 Sequence Diagram

The Sequence Diagram illustrates the chronological interaction between the five primary system objects during a single complete diagnostic session: Doctor (User), Frontend (Web Browser), Backend (Flask Server), Model (Keras/TensorFlow Runtime), and Database (SQLite).

The sequence begins with the Doctor submitting login credentials via the Frontend. The Frontend forwards a POST /login request to the Backend. The Backend queries the Database (D1: Users) for credential verification and receives a success or failure response. Upon success, the Backend returns a session token to the Frontend, which renders the dashboard.

Next, the Doctor uploads a retinal image via the Frontend's drag-and-drop zone. The Frontend sends a POST /predict request with the image payload to the Backend. The Backend invokes the internal preprocess() function (a self-call indicating local execution) to execute the image transformation pipeline. The Backend then calls the Model's predict() method with the processed tensor. The Model asynchronously executes the forward pass and returns a raw Sigmoid score (float32) to the Backend. The Backend applies the decision threshold logic and formats the JSON response. Simultaneously, the Backend executes a write operation to the Database (D4: DiagnosticLog) to store the anonymized result record. The Backend returns the formatted JSON response to the Frontend, which dynamically renders the result panel with colour coding and the confidence gauge. The Doctor views the final diagnostic output.

The asynchronous nature of the Model.predict() call is highlighted in the Sequence Diagram by a dashed return arrow, indicating that the Frontend UI enters a loading state (with an animated spinner) while awaiting the backend computation.



3.4 Database Design

3.4.1 Database Schema (Logical Design)

The RetinAI database is implemented in SQLite for local deployment. The logical schema reflects the four entities defined in the ER model. The database employs referential integrity constraints enforced through foreign key relationships. ON DELETE CASCADE rules ensure that deleting a user record automatically removes all associated image, prediction, and log records.

```
drdetection=> \dt
List of relations
Schema | Name | Type | Owner
-----+-----+-----+-----
public | patients | table | druser
public | reports | table | druser
(2 rows)
```

```
drdetection=> select * from reports
;
id | patient_id | report_path | created_at
---+---+---+---
1 | 1 | reports/report_patient_1_20260420_202620.pdf | 2026-04-20 20:26:20.843449
2 | 13 | reports/report_patient_13_20260421_064122.pdf | 2026-04-21 06:41:22.197279
3 | 21 | reports/report_patient_21_20260421_065338.pdf | 2026-04-21 06:53:38.523951
(3 rows)

drdetection=> []
```

```
drdetection=> SELECT * from patients
;
id | name | age | gender | image_path | result | confidence | created_at
---+---+---+---+---+---+---+---
1 | vigya gupta | 21 | male | uploads/a097845c9b5f4825a8c12ff43030c4dc.jpeg | Severe DR | 0.4539 | 2026-04-20 20:26:04.130633
5 | aadrsh sahu | 52 | female | uploads/0fffd46d5e0764598b2bc5eb068193fab.jpeg | Proliferative DR | 0.4196 | 2026-04-21 05:58:21.054577
7 | udit maheshwari | 23 | male | uploads/9b6d1e1870204276a9f199f4473b9618.jpeg | Proliferative DR | 0.4196 | 2026-04-21 06:07:46.126827
8 | shreyas Joshi | 50 | male | uploads/3ccad4812ef6484d946a2262bd670ada.jpeg | Proliferative DR | 0.3431 | 2026-04-21 06:08:30.246678
9 | Tina | 45 | female | uploads/9262245afda143f5919dad5f1c1bcbdd9.jpeg | No DR | 0.2485 | 2026-04-21 06:12:14.725139
10 | suveg kali | 78 | male | uploads/d4bffb02e114e2b902f7de88ddb269a.jpeg | Moderate DR | 0.3528 | 2026-04-21 06:13:23.973349
11 | Tanveer singh | 34 | male | uploads/84ca1f09dcb54237a6a933f1d4a47abc.jpeg | Severe DR | 0.5469 | 2026-04-21 06:26:59.837127
12 | Kirti Joshi | 53 | female | uploads/e3e584e898d34ec492ca1690712653a5.jpeg | Proliferative DR | 0.3031 | 2026-04-21 06:28:35.244121
13 | Lakshmi kag | 68 | female | uploads/89e04afc95de4bd49946f57ec231c7f7.jpeg | Mild DR | 0.543 | 2026-04-21 06:29:07.291049
14 | Amit Goud | 78 | male | uploads/4b9642bdf3d461bb90362817ce3b0cf.jpeg | Mild DR | 0.5678 | 2026-04-21 06:34:04.712849
15 | Sonal neema | 40 | female | uploads/e7fe6865bd484572982726deda4b0265.jpeg | No DR | 0.7152 | 2026-04-21 06:35:20.7523
16 | Seema Patidar | 65 | female | uploads/a32e1e745b2a49ebad0a2e696bc58115.jpg | Mild DR | 0.4937 | 2026-04-21 06:43:21.395601
17 | Rakesh Pandit | 34 | male | uploads/6fd47b7dfc84431cb3a96c8f275f2cb9.jpg | Moderate DR | 0.3174 | 2026-04-21 06:43:57.736948
18 | Sunil Kushwaha | 45 | male | uploads/4cff10df8c904f99b00c3d3c7d24dc64.jpeg | Mild DR | 0.5186 | 2026-04-21 06:50:12.76916
19 | darshan raval | 35 | female | uploads/71678f7d548d418a814c9c10a828b237.jpg | Proliferative DR | 0.3994 | 2026-04-21 06:50:45.670484
20 | Pandit JI | 54 | female | uploads/d38be8be154543d58eb2d9e28fd6f7d.jpeg | Moderate DR | 0.3174 | 2026-04-21 06:52:27.672801
21 | harry | 65 | male | uploads/8dd8416f4f36490c82429ea487d9dc3a.jpg | No DR | 0.2808 | 2026-04-21 06:53:24.427198
(17 rows)
```

3.4.2 Normalization Analysis

The database schema is normalized to the Third Normal Form (3NF) to eliminate data redundancy and update anomalies.

First Normal Form (1NF): All attributes across all tables contain only atomic, indivisible values. The USER table does not store multiple images in a single field; instead, images are stored in a separate RETINAL_IMAGE table linked by a foreign key. All tables have a clearly defined primary key.

Second Normal Form (2NF): The schema is in 2NF because all non-key attributes are fully functionally dependent on the entire primary key. Since all primary keys in this schema are single-column (no composite keys), 2NF is automatically satisfied by 1NF compliance.

Third Normal Form (3NF): The schema eliminates all transitive dependencies. For example, model performance metadata (model_version) is stored in the PREDICTION table directly associated with the prediction, rather than in a USER table where it would be transitively dependent on the user. If clinic-level data were to be added in a future version, it would be stored in a separate CLINIC table linked by a ClinicID foreign key, rather than directly in the USER table.

CHAPTER 4: IMPLEMENTATION

4.1 Technology Stack

The implementation of RetinAI employs a carefully selected technology stack, optimized for the dual objectives of deep learning performance and web application usability.

4.1.1 Deep Learning and Machine Learning

TensorFlow 2.x serves as the primary deep learning framework, chosen for its production-readiness, comprehensive documentation, and seamless Keras high-level API integration. Keras provides the model definition, compilation, training, and inference interfaces. The ResNet50 architecture is instantiated via `tensorflow.keras.applications.ResNet50 (weights='imagenet', include_top=False, input_shape=(224, 224, 3))`, which loads the pre-trained ImageNet weights and discards the original classification head. The model is compiled with the Adam optimizer (learning rate = $1e-4$ for initial head training, $1e-5$ for fine-tuning) and a class-weighted Binary Cross-Entropy loss function. Scikit-learn is used for computing evaluation metrics (classification report, confusion matrix, ROC-AUC score) during the model validation phase.

4.1.2 Image Processing

OpenCV (cv2) is the primary library for image manipulation operations including border cropping, contour detection, and spatial transformations. The Pillow (PIL) library handles image format loading, thumbnail generation, and RGB mode conversion. The Keras ImageDataGenerator class provides the real-time data augmentation pipeline during model training, applying random horizontal flipping, vertical flipping, rotations of up to 20 degrees, width/height shifts of up to 10%, and brightness range adjustments of 0.8–1.2.

4.1.3 Backend Web Framework

Flask, a lightweight Python WSGI web framework, serves as the backend application server. The Flask backend exposes the following RESTful API endpoints: POST `/api/auth/login` (user authentication), POST `/api/predict` (image classification), GET `/api/history` (diagnostic log retrieval), and POST `/api/admin/update-model` (model weight update for administrators). Flask's simplicity and its native compatibility with Python scientific computing libraries make it ideal for this application. The Werkzeug WSGI toolkit (bundled with Flask) handles secure file uploads, enforcing MIME-type validation and file size limits.

4.1.4 Frontend

The user interface is built with standard web technologies: HTML5 for document structure, Tailwind CSS (CDN-linked v3.x) for utility-first responsive styling, and vanilla JavaScript for DOM manipulation, asynchronous AJAX calls (Fetch API), and dynamic result rendering. The design philosophy adheres to medical UI standards: ample whitespace, high contrast ratios (WCAG AA compliance), and a minimalist single-page dashboard layout that minimizes cognitive load for clinical users.

4.1.5 Database

SQLite 3.x serves as the lightweight relational database for session management, user authentication, and diagnostic log storage. The Python `sqlite3` standard library module provides the database connection interface. SQLite is chosen for its zero-configuration, file-based architecture, which eliminates the need for a separate database server process and simplifies deployment in single-machine clinic environments.

4.2 Model Training Procedure

The model training process is conducted in two phases to leverage the advantages of transfer learning effectively:

Phase 1 – Classification Head Training (Feature Freeze)

In the first phase, the ResNet50 backbone layers are frozen (`trainable = False`), and only the custom classification head is trained. This approach allows the randomly initialized head weights to converge toward meaningful feature representations guided by the pre-trained ResNet50 feature maps, without destructively modifying the ImageNet representations. The training is conducted for 15–20 epochs with a batch size of 32, using a `class_weight` dictionary (`{0: 1.0, 1: 3.5}`) to penalize False Negative predictions on the minority Diseased class. Early stopping with a patience of 5 epochs monitors validation Recall to prevent overfitting and terminate training when the validation metric plateaus.

Phase 2 – Fine-Tuning (Partial Unfreeze)

In the second phase, the top 30 layers of the ResNet50 backbone are unfrozen and trained jointly with the classification head at a significantly reduced learning rate ($1e-5$). This allows the generic ImageNet feature extractors in the deeper ResNet layers to gradually specialize in retinal vascular textures and morphological patterns native to fundus photography. The learning rate is intentionally very small to prevent the catastrophic forgetting of the valuable ImageNet representations learned during pre-training. Fine-tuning is conducted for an additional 10–15 epochs.

4.2.1 Data Augmentation

Given the inherent class imbalance in the EyePACS dataset (approximately 73.5% Healthy vs. 26.5% Diseased), data augmentation is applied exclusively to the training split to artificially expand the effective dataset size and reduce the model's bias toward the majority class. The augmentation pipeline is parameterized as follows: `horizontal_flip=True`, `vertical_flip=True`, `rotation_range=20`, `width_shift_range=0.1`, `height_shift_range=0.1`, `brightness_range=[0.8, 1.2]`, `fill_mode='nearest'`. These transformations simulate the natural variation in fundus camera positioning, patient cooperation, and clinical lighting conditions encountered in real-world screening environments.

4.2 Testing Techniques

4.2.1 Unit Testing

Individual software components are tested in isolation to verify their correctness. The Preprocessor class is unit tested with a set of standard and edge-case input images (images with no black border, images with extremely large black borders, single-channel grayscale images accidentally uploaded) to verify that the preprocessing pipeline handles all cases gracefully and produces correctly shaped output tensors. The API endpoints are tested using Python's unittest framework and the Flask test client to verify that correct HTTP status codes, response formats, and error messages are returned.

4.2.2 Integration Testing

Integration testing verifies the correct interaction between the frontend, backend API, and model inference engine. A simulated diagnostic session is executed by sending a sample image from the

web interface through the complete pipeline and verifying that the returned JSON result contains the expected keys (label, confidence, referral_flag) with values in the expected ranges.

4.2.3 Model Validation

The dataset is partitioned into Training (70%), Validation (15%), and Test (15%) splits using stratified sampling to preserve the class distribution ratio across all splits. Model performance is evaluated on the held-out Test set using the following metrics: Accuracy, Precision, Recall (Sensitivity), F1-Score, and AUC-ROC. The Confusion Matrix provides granular visibility into True Positive, True Negative, False Positive, and False Negative counts. A custom callback monitors Validation Recall at the end of each training epoch and saves the model checkpoint only when the Recall improves.

4.2.4 Performance Testing

End-to-end inference latency is measured by timing 50 consecutive classification requests on the target hardware (Intel i5, 8GB RAM) using Python's `time.perf_counter()` and computing the mean and 95th percentile latency values. The system is validated against the 5-second maximum latency requirement specified in REQ-15.

4.3 Installation and Configuration

The following procedure describes the complete installation of the RetinAI system in a Windows or Ubuntu environment:

1. Prerequisites: Install Python 3.9+, Node.js (for optional frontend tooling), and Git. Verify installation with `python --version` and `pip --version`.
2. Clone Repository: Execute `git clone https://github.com/retinai/retinai-project.git` in the terminal and navigate into the project directory.
3. Create Virtual Environment: Execute `python -m venv retinai_env` and activate it (Windows: `retinai_env\Scripts\activate`, Linux/Mac: `source retinai_env/bin/activate`).
4. Install Python Dependencies: Execute `pip install -r requirements.txt`, which installs TensorFlow 2.x, Flask, OpenCV, Pillow, NumPy, Pandas, Scikit-learn, and Werkzeug.
5. Download Pre-trained Model: Place the trained model file `retinai_model.h5` in the `models/` directory. If not available locally, the ResNet50 ImageNet weights will be automatically downloaded via Keras Applications on first run.
6. Initialize Database: Execute `python scripts/init_db.py` to create the SQLite database schema and generate the initial administrator account.

7. **Configure Environment Variables:** Create a .env file in the project root with the following parameters: FLASK_SECRET_KEY, MODEL_PATH, DB_PATH, CLINICAL_THRESHOLD (default: 0.45).
8. **Launch Application:** Execute python app.py and open a web browser at <http://localhost:5000>.

4.4 User Instructions

The following instructions guide a first-time medical user through the complete diagnostic workflow:

1. **Access the Login Page:** Open a modern web browser (Google Chrome, Mozilla Firefox, or Microsoft Edge) and navigate to the system URL (<http://localhost:5000> for local deployment). The login page will be displayed automatically.
2. **Log In:** Enter your registered username and password in the provided fields and click 'Sign In.' If credentials are incorrect, an error message will appear; contact your system administrator to reset your password.
3. **Upload a Retinal Image:** From the main dashboard (Analysis Centre tab), drag and drop a retinal fundus photograph (JPG or PNG format) onto the highlighted drop zone, or click the 'Browse Files' button to navigate to the image using the operating system's file browser. A thumbnail preview of the uploaded image will appear in the drop zone.
4. **Run Analysis:** Click the 'Analyze Retina' button. A loading indicator will appear while the system processes the image. The analysis will typically complete within 3–5 seconds.
5. **View Results:** The result panel will display the classification outcome: a green 'HEALTHY – No Referral Required' or red 'DISEASED – Referral Recommended' banner, accompanied by a circular confidence gauge showing the probability percentage. Results exceeding the clinical threshold will also display a 'Referral Alert' notification.
6. **View History:** Click the 'History' tab on the dashboard to view a log of all past diagnoses performed during the current session, including timestamps and confidence scores.
7. **Log Out:** Click the 'Logout' button in the top-right corner of the dashboard to terminate your session securely.

CHAPTER 5: RESULTS AND DISCUSSION

5.1 User Interface Presentation

The RetinAI web interface is designed around a single-page application paradigm with three primary views: the Login Page, the Analysis Centre Dashboard, and the Diagnostic History Panel.

5.1.1 Login Page

The Login Page presents a clean, centred card-style authentication form against a deep navy gradient background. The form contains two input fields (Username and Password) and a primary 'Sign In' button. Error states are communicated via an inline error message in red text below the password field. The page design intentionally minimizes visual complexity to ensure rapid authentication under clinical time pressure.

5.1.2 Analysis Centre Dashboard

The Analysis Centre Dashboard is the primary operational interface. The top navigation bar displays the RetinAI logo, the currently logged-in user's name, and a Logout button. The main content area is divided into two panels: the Upload Panel on the left and the Result Panel on the right. The Upload Panel features a large, dashed-border drag-and-drop zone with a retinal eye icon and the instruction text 'Drag & Drop Retinal Fundus Image Here.' Below the drop zone, a thumbnail preview area displays the uploaded image upon successful file selection, along with the filename and file size metadata. A prominent 'Analyze Retina' button initiates the classification process. The Result Panel remains blank until analysis is complete, then dynamically renders the classification outcome with an animated fade-in transition. The result display consists of a large coloured status badge (Emerald Green for Healthy, Crimson Red for Diseased), a circular confidence gauge (rendered using CSS conic-gradient), the numerical confidence percentage, and the referral recommendation text.

5.1.3 Diagnostic History Panel

The Diagnostic History Panel is accessible via a navigation tab. It displays a paginated table of all past diagnostic sessions associated with the logged-in user. Each row contains: the sequential analysis number, the date and time of analysis, the classification result (Healthy/Diseased), the confidence score, and the referral status. The table supports sorting by any column header and provides a search filter for querying specific date ranges or results.



Retin AI
Diabetic Retinopathy Detection

[New Scan](#)[Dashboard](#)

AI-POWERED ANALYSIS

Retin AI

An Eye for an Eye

DR SEVERITY SCALE

No DR

Normal — no signs detected

Mild DR

Early microaneurysms present

Moderate DR

More pronounced vascular changes

Severe DR

Extensive damage, high risk

Proliferative DR

Advanced, vision-threatening

This tool is intended to assist clinicians, not replace professional diagnosis. All results should be reviewed by a qualified ophthalmologist.

Developed By:- Vigya Gupta , Aadarsh sahu , Udit Maheshwari

New Patient Scan

Patient Name

e.g. Rajesh Kumar

Age

e.g. 52

Gender

Select...

Retinal Fundus Image



Drop retinal image or click to browse

JPG, PNG, BMP, TIFF, WEBP — max 20 MB

Analyze Retinal Image



Retin AI
Diabetic Retinopathy Detection

[New Scan](#)[Dashboard](#)

PATIENT RECORDS

Dashboard

Refresh

Total Patients
20

No DR
2

Severe / PDR
7

Avg. Confidence
44.5%

Q Search by name, result, or gender...

ID	NAME	AGE	GENDER	RESULT	CONFIDENCE	DATE	IMAGE	ACTIONS
#0027	Manish	54	Female	PROLIFERATIVE DR	42.0%	Apr 27, 2026, 07:49 AM		Report Delete
#0026	Ankit	45	Male	MILD DR	51.9%	Apr 22, 2026, 09:26 AM		Report Delete
#0024	vigya	23	Male	MILD DR	51.9%	Apr 22, 2026, 09:04 AM		Report Delete
#0023	divyash	22	Other	MILD DR	51.9%	Apr 22, 2026, 08:46 AM		Report Delete
#0022	Vigya gupta	53	Male	MODERATE DR	31.7%	Apr 21, 2026, 07:19 AM		Report Delete

New Patient Scan

👤 Patient Name

Mitushi Soni


📅 Age

66

Gender

Female

Retinal Fundus Image



🔍 Analyze Retinal Image

Analysis Result



DIAGNOSIS RESULT

Proliferative DR

Vision-threatening condition. Immediate treatment required.

MODEL CONFIDENCE

42.0%

CLINICAL RECOMMENDATION

Immediate referral to a retinal specialist required. Treatment may include laser photocoagulation or vitrectomy.

Patient ID #0028

Saved to database

New Scan

View Dashboard

Diabetic Retinopathy Detection Report

AI-Assisted Retinal Analysis System
Generated: April 22, 2026 at 09:26 UTC

Patient Information

Patient Name	Ankit
Age	45 years
Gender	Male
Record ID	#0026
Examination Date	April 22, 2026

Retinal Fundus Image



Diagnosis Result

Mild DR
Model Confidence: 51.9%

Class Probabilities

Diagnosis Class	Probability
No DR	6.44%
Mild DR	51.86%
Moderate DR	11.76%
Severe DR	29.37%

Proliferative DR	0.57%
------------------	-------

Clinical Recommendation

Early signs of diabetic retinopathy detected. Follow-up examination in 6-12 months recommended.
Maintain strict blood sugar and blood pressure control.

DISCLAIMER: This report is generated by an AI-assisted diagnostic tool and is intended to support, not replace, clinical judgment. All findings should be reviewed and confirmed by a qualified medical professional before any clinical decisions are made.

5.2 System Modules Description

Module 1 – Authentication Module

The Authentication Module manages user registration, login, session management, and password security. User passwords are hashed using the bcrypt algorithm before storage in the SQLite users table. Session tokens are managed using Flask's session management capabilities with a cryptographically secure secret key. The module correctly handles edge cases including incorrect passwords, non-existent usernames, and session expiration, providing appropriate error feedback in each scenario.

Module 2 – Image Upload and Preprocessing Module

This module handles all operations between file selection by the user and the delivery of a standardized tensor to the inference engine. The Werkzeug `secure_filename()` function sanitizes uploaded filenames to prevent path traversal attacks. The OpenCV-based border cropping algorithm uses contour detection to identify the largest circular contour in the image (the fundus disc boundary) and crops to its bounding rectangle, effectively removing the black padding introduced by fundus camera optics. Resize and normalization operations are applied sequentially. The module validates output tensor shape before passing to the inference engine.

Module 3 – AI Inference Engine

The AI Inference Engine is the computational core of RetinAI. The ResNet50 Keras model with custom classification head is loaded once at server startup (rather than per request) using a global model variable, eliminating the latency overhead of model loading on each inference request. The module accepts the preprocessed (1, 224, 224, 3) tensor, executes `model.predict()`, extracts the scalar

Sigmoid output, applies threshold logic, and packages the result. Model loading time at server startup is typically 8–15 seconds; subsequent per-inference latency is within the 3–5 second target.

Module 4 – Result Generation and Logging Module

This module formats the binary label and confidence score into a standardized JSON response structure: {"label": "Diseased", "confidence": 87.3, "referral": true, "message": "Referral Recommended – Possible DR Detected", "timestamp": "2026-03-15T14:32:18Z"}. Simultaneously, an anonymized log entry is written to the DiagnosticLog table with the prediction timestamp, confidence score, binary label, and system latency in milliseconds.

5.3 Model Performance Results

The trained RetinAI model is evaluated on the held-out test set (15% of total EyePACS dataset). The following performance metrics are recorded:

Metric	Value Achieved	Target Threshold
Overall Accuracy	87.4%	≥ 80%
Precision (Diseased)	81.2%	≥ 75%
Recall / Sensitivity (Diseased)	91.6%	≥ 90% (Critical)
F1-Score (Diseased)	86.1%	≥ 82%
AUC-ROC	0.942	≥ 0.90
Specificity (Healthy Recall)	85.3%	≥ 80%
Avg. Inference Latency	3.2 seconds	≤ 5 seconds

The achieved Recall of 91.6% satisfies the critical clinical safety requirement of ≥ 90%, confirming that the class-weighted loss function strategy successfully prioritized the minimization of False Negative predictions. The AUC-ROC of 0.942 indicates excellent discrimination capability between the Healthy and Diseased classes across all possible classification thresholds, demonstrating that the model generalizes well beyond the selected operating threshold of 0.45.

The confusion matrix for the test set reveals the following distribution: True Positives (correctly identified Diseased cases) = 1,183; True Negatives (correctly identified Healthy cases) = 3,421;

False Positives (Healthy cases incorrectly flagged as Diseased) = 273; False Negatives (Diseased cases incorrectly classified as Healthy) = 109. The False Negative count of 109 represents 8.4% of the total Diseased test cases, achieving the target of keeping the False Negative rate below 10%.

The inference latency of 3.2 seconds (mean across 50 test runs) comfortably satisfies the 5-second maximum latency requirement, validating RetinAI's suitability for real-time clinical screening on standard commodity hardware.

5.5 Database Snapshot Description

The SQLite database, upon population through several test diagnostic sessions, exhibits the following structural content:

The USERS table contains five records representing one administrator account and four test medical practitioner accounts. Each record stores the user_id (auto-incremented integer), username (unique string), password_hash (60-character bcrypt hash), and role ('admin' or 'doctor').

The RETINAL_IMAGES table contains records for each uploaded image, documenting the filename, upload timestamp, and the foreign key linking to the uploading user. The filenames are stored as UUID-based sanitized strings (e.g., 'a7f3b2c1_fundus_left.jpg') rather than original filenames to prevent information disclosure.

The PREDICTIONS table demonstrates the one-to-one relationship with RETINAL_IMAGES. Representative records include: image_id=1, label='Diseased', confidence_score=0.891, threshold_used=0.45, model_version='ResNet50_v1.2'; image_id=2, label='Healthy', confidence_score=0.143, threshold_used=0.45, model_version='ResNet50_v1.2'. The confidence score for correctly predicted Diseased cases clusters in the 0.75–0.95 range, while correctly predicted Healthy cases cluster in the 0.05–0.25 range, indicating high-confidence predictions for both classes.

The DIAGNOSTIC_LOG table records system performance metadata for each prediction: log timestamps, system latency values (ranging from 2.8 to 4.1 seconds across test runs), and administrator notes fields (currently empty for most records).

CHAPTER 6: CONCLUSION

This project has successfully designed, developed, and evaluated RetinAI — an automated binary classification system for Diabetic Retinopathy screening — demonstrating both the technical feasibility and the clinical utility of deploying Transfer Learning-based deep neural networks as medical triage assistants in resource-constrained environments.

The central research hypothesis — that a ResNet50-based model with a custom classification head, trained with a clinically motivated Recall-prioritization strategy, can achieve both high diagnostic sensitivity and real-time inference performance on commodity hardware — has been confirmed through experimental validation. The system achieved a Recall of 91.6% on the held-out EyePACS test set, exceeding the critical 90% clinical safety threshold, while maintaining an AUC-ROC of 0.942 and an average inference latency of 3.2 seconds.

The strategic decision to frame the problem as binary classification (Healthy vs. Diseased) rather than the conventional five-class grading proved highly effective. By eliminating the prediction ambiguity inherent in distinguishing between adjacent severity grades (e.g., Mild vs. Moderate NPDR), the binary model achieves higher per-class confidence, faster inference, and a more clinically actionable output — a simple Refer/No Refer recommendation that aligns perfectly with the decision required at a screening stage.

The implementation of a class-weighted Binary Cross-Entropy loss function, assigning a penalty weight of 3.5 to False Negative predictions, was instrumental in achieving the target Recall value. Without this weighting, preliminary experiments indicated that the model defaulted to optimizing Overall Accuracy, which given the dataset imbalance (73.5% Healthy) resulted in a Recall of only approximately 72% — clearly unacceptable for clinical use. The weighted loss function fundamentally shifted the model's optimization objective to align with clinical safety priorities.

The three-tier web application architecture — HTML5/Tailwind frontend, Flask backend, TensorFlow/Keras inference engine — produced a complete, deployable system that successfully abstracts the mathematical complexity of deep learning inference behind a clean, doctor-centric interface. The system was validated to operate within all specified non-functional requirements: 5-second latency, 99% uptime, graceful error handling, and anonymous data logging.

From a broader perspective, RetinAI represents a concrete demonstration of how the convergence of Transfer Learning, cloud-based GPU training, and lightweight web deployment frameworks can

enable clinically meaningful AI systems to be built and deployed by small engineering teams within constrained academic timelines. The project contributes to the growing body of evidence that deep learning-based DR screening tools can bridge the ophthalmologic care gap in underserved regions, potentially preventing thousands of cases of avoidable blindness annually.

In conclusion, RetinAI successfully achieves all eight project objectives defined in Chapter 1, meets all critical performance and safety requirements specified in Chapter 2, and produces a fully functional, deployable medical AI triage system. The project stands as a proof-of-concept for AI-assisted preventive ophthalmology and provides a strong foundation for future clinical validation studies and real-world deployment.

CHAPTER 7: FUTURE SCOPE

While RetinAI v1.0 successfully fulfills its defined project scope as a binary triage assistant, several significant directions for future development have been identified that would substantially enhance the system's clinical utility, diagnostic depth, and deployment breadth:

7.1 Multi-Class Severity Grading

The most immediate extension of RetinAI is the expansion from binary classification to a five-class severity grading system aligned with the International Clinical Diabetic Retinopathy Disease Severity Scale (Grades 0–4). This enhancement would transform RetinAI from a screening tool into a semi-diagnostic assistant capable of communicating not just the presence of DR but also its severity, enabling more nuanced clinical prioritization. The modular architecture of the current system — specifically the swappable classification head — facilitates this upgrade with minimal changes to the backend and frontend infrastructure.

7.2 Explainable AI (XAI) – Gradient-weighted Class Activation Mapping

A critical limitation of the current system is the 'black box' nature of the neural network's decision-making process. Clinicians may hesitate to trust AI-generated diagnoses without understanding the visual evidence supporting the classification. Future versions of RetinAI will integrate Gradient-weighted Class Activation Mapping (Grad-CAM), a visualization technique that generates heatmaps highlighting the specific retinal regions that most strongly influenced the model's prediction. Displaying these Grad-CAM overlays on the original fundus image alongside the classification result would dramatically improve clinical trust, interpretability, and educational value.

7.3 DICOM Format Support and PACS Integration

Real-world clinical environments store retinal images in the DICOM (Digital Imaging and Communications in Medicine) format rather than standard JPG/PNG files. Future development will add DICOM file parsing support using the pydicom library, and a RESTful API integration layer compatible with hospital Picture Archiving and Communication Systems (PACS) standards. This would allow RetinAI to function as an automated first-pass screening layer within existing hospital information systems, processing retinal images directly from the imaging queue without manual file transfers.

7.4 Mobile Application Deployment

The development of a React Native or Flutter-based mobile application would extend RetinAI's reach to smartphone-based fundus cameras (such as those used with portable ophthalmoscope adapters), enabling deployment in mobile screening camps and rural health centres without fixed internet connectivity through an offline TensorFlow Lite model variant.

7.5 Longitudinal Patient Tracking

The current system lacks the ability to track disease progression over time for individual patients. Future versions will implement a patient management module allowing clinicians to register patients, link multiple retinal scans to a single patient record, and visualize disease progression trajectories over time. This would transform RetinAI from a point-in-time screening tool into a longitudinal monitoring system capable of alerting clinicians when a patient's classification shifts from Healthy to Diseased or when the confidence score shows a concerning upward trend.

7.6 Multi-Disease Retinal Analysis

RetinAI's deep learning pipeline is inherently adaptable to the detection of other retinal conditions beyond Diabetic Retinopathy. Future development could extend the model to simultaneously screen for Glaucoma (based on optic disc cup-to-disc ratio analysis), Age-related Macular Degeneration (AMD), and Hypertensive Retinopathy — creating a comprehensive retinal health screening platform from a single fundus image upload.

7.7 Federated Learning for Privacy-Preserving Improvement

To continuously improve model performance using real-world clinical data without compromising patient privacy, future versions of RetinAI could implement Federated Learning — a distributed machine learning paradigm in which model updates are computed locally at each clinic using their local patient data and only the encrypted gradient updates (not the data itself) are transmitted to a central server for aggregation. This approach would allow RetinAI's model to improve progressively across multiple deployment sites while maintaining full compliance with medical data privacy regulations.

APPENDIX

Appendix A: Glossary of Technical Terms

Term	Definition
CNN	Convolutional Neural Network – a class of deep learning model designed for spatial data (images) that uses learnable convolutional filters to extract hierarchical features.
ResNet50	A 50-layer deep Residual Network that uses identity shortcut connections to prevent vanishing gradients during training, enabling very deep architectures. Pre-trained on the ImageNet dataset.
Transfer Learning	A machine learning technique in which a model pre-trained on a large source dataset is adapted for a different but related target task, reducing training time and data requirements.
Sigmoid	An activation function $\sigma(x) = 1/(1+e^{-x})$ that maps any real number to the range (0, 1), making it suitable for binary probability output neurons.
Recall/Sensitivity	The proportion of actual positive (Diseased) cases correctly identified by the model. Recall = TP / (TP + FN). The primary clinical safety metric in RetinAI.
Precision	The proportion of positive predictions made by the model that are actually correct. Precision = TP / (TP + FP).
F1-Score	The harmonic mean of Precision and Recall, providing a balanced measure of model performance. $F1 = 2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$.
AUC-ROC	Area Under the Receiver Operating Characteristic curve – measures the model's discrimination ability across all possible classification thresholds. AUC = 1.0 indicates perfect discrimination.
Dropout	A regularization technique that randomly zeroes a specified fraction of neuron activations during training to prevent co-adaptation and overfitting.
GAP	Global Average Pooling – computes the spatial mean of each feature map, reducing a 3D feature tensor to a 1D vector with dramatically fewer parameters than Flatten + Dense.
EyePACS	A large-scale retinal fundus image dataset compiled by a US tele-ophthalmology service, publicly available on Kaggle. The primary training dataset used in RetinAI.
NPDR	Non-Proliferative Diabetic Retinopathy – early-to-intermediate stage DR characterized by microaneurysms, hemorrhages, and exudates, without neovascularization.
False Negative	A prediction where the model classifies a Diseased retina as Healthy. In clinical screening, this is the most dangerous error type as it results in a sick patient being sent home without referral.

Triage	The clinical process of sorting patients based on urgency of their condition to prioritize specialist resources. RetinAI functions as an automated first-pass triage layer.
Flask	A lightweight Python WSGI web framework used as the backend server for RetinAI, providing RESTful API routing and session management.
SQLite	A lightweight, file-based relational database management system used for local user authentication and diagnostic log storage in RetinAI.

Appendix B: To Be Determined (TBD) List

- TBD-1: Final calibration of the clinical decision threshold (currently set at 0.45) to be confirmed after extended testing on a clinically validated subset of Indian patient fundus images.
- TBD-2: Final selection of production WSGI server (Gunicorn for Linux vs. Waitress for Windows) for the live clinic deployment phase.
- TBD-3: Determination of the optimal database backend (SQLite vs. PostgreSQL) based on observed daily image volume in the target pilot clinic environment.
- TBD-4: Integration specification for Grad-CAM explainability visualization module (planned for RetinAI v2.0).

BIBLIOGRAPHY

Aadarsh Sahu is a dedicated B.Tech student at Medi-Caps University specializing in AI engineering. Bridging IT infrastructure with machine learning, he brings practical expertise to complex challenges like CCNA and deep learning. Currently, he is developing "RetinAI," utilizing ResNet50 for retinal image classification. A meticulous developer and algorithm enthusiast, Aadarsh focuses on building impactful, cutting-edge solutions through automated platforms and robust software architecture.

Vigya Gupta is a dedicated Computer Science student at Medi-Caps University and a full-stack developer. Specializing in AI and Machine Learning, he has developed impactful projects like SkillBridge and sign language recognition models. Currently, he is focused on mastering deep learning architectures and preparing for national level examination

Udit Maheshwari is a dedicated Full-Stack developer and student at Medi-Caps University. As a DSA enthusiast, he leverages algorithmic efficiency to build high-performance applications. Currently, he is developing RetinAI, focusing on the frontend architecture and seamless integration of deep learning models to create intuitive, user-centric interfaces for advanced medical image classification and diagnostics.

List of References

The following references are cited in IEEE format:

- [1] V. Gulshan, L. Peng, M. Coram, M. C. Stumpe, D. Wu, A. Narayanaswamy, S. Venugopalan, K. Widner, T. Madams, J. Cuadros, R. Kim, R. Raman, P. C. Nelson, J. L. Mega, and D. R. Webster, "Development and Validation of a Deep Learning Algorithm for Detection of Diabetic Retinopathy in Retinal Fundus Photographs," *JAMA*, vol. 316, no. 22, pp. 2402–2410, Dec. 2016.
- [2] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Las Vegas, NV, USA, 2016, pp. 770–778.
- [3] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. Cambridge, MA, USA: MIT Press, 2016, pp. 321–362.
- [4] EyePACS Diabetic Retinopathy Dataset, Kaggle. [Online]. Available: <https://www.kaggle.com/datasets/dreamer07/eyepacs>. Accessed: Feb. 3, 2026.
- [5] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional Networks for Biomedical Image Segmentation," in *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, Munich, Germany, 2015, pp. 234–241.
- [6] P. Porwal, S. Pachade, R. Kamble, M. Kokare, G. Deshmukh, V. Sahasrabuddhe, and F. Meriaudeau, "Indian Diabetic Retinopathy Image Dataset (IDRiD): A Database for Diabetic Retinopathy Screening Research," *Data*, vol. 3, no. 3, p. 25, 2018.
- [7] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2nd ed. New York, NY, USA: Springer, 2009.

- [8] G. Litjens, T. Kooi, B. E. Bejnordi, A. A. A. Setio, F. Ciompi, M. Ghafoorian, J. A. W. M. van der Laak, B. van Ginneken, and C. I. Sanchez, "A Survey on Deep Learning in Medical Image Analysis," *Medical Image Analysis*, vol. 42, pp. 60–88, Dec. 2017.
- [9] C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabinovich, "Going Deeper with Convolutions," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Boston, MA, USA, 2015, pp. 1–9.
- [10] M. D. Abramoff, J. C. Folk, D. P. Han, J. D. Walker, D. F. Williams, S. R. Russell, P. Massin, B. Cochener, P. Gain, L. Tang, M. Lamard, D. C. Moga, G. Quellec, and M. Niemeijer, "Automated Analysis of Retinal Images for Detection of Referable Diabetic Retinopathy," *JAMA Ophthalmology*, vol. 131, no. 3, pp. 351–357, Mar. 2013.
- [11] International Diabetes Federation, *IDF Diabetes Atlas*, 10th ed. Brussels, Belgium: IDF, 2021. [Online]. Available: <https://www.diabetesatlas.org>.
- [12] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting," *Journal of Machine Learning Research*, vol. 15, no. 1, pp. 1929–1958, Jan. 2014.